

Project

April 19, 2025

1 PHYS 310 FINAL PROJECT

1.0.1 Load the Dataset

Dataset is obtained from the paper:

“Prediction and Classification of Formation Energies of Binary Compounds by Machine Learning” https://pubs.acs.org/doi/10.1021/acsomega.1c01517?goto=supporting-info#_i28

```
[14]: import pandas as pd  
file_path = "ao1c01517_si_003.xlsx"  
df = pd.read_excel(file_path, skiprows = 51)  
df
```

```
[14]:    Substance   H(J/mol-atomes)   S0 J/cal/mol   BDE(kj/mol)   SA(J/cal/deg)  \  
0      Li2C2     -14850.0000     5.770000       0.0          29.1  
1      Li3N     -41200.0000     4.681075       0.0          29.1  
2      Li2O     -199300.0000     5.298567      340.5         29.1  
3      Li2O2    -158475.0000     5.770000      340.5         29.1  
4      LiF      -307500.0000     5.770000      577.0         29.1  
..      ...        ...          ...          ...          ...  
178     BaF2     -402933.3333     5.298567      580.0         62.4  
179     BaS      -231800.0000     5.770000       0.0          62.4  
180     BaCl2    -286200.0000     5.298567      439.3         62.4  
181     BaBr2    -252566.6667     5.298567      359.9         62.4  
182     Ba2Sn    -125533.3333     5.298567       0.0          62.4  
  
      SB(J/cal/mol)   A_IP(ev)   A_EA(ev)   A_EN(ev)   A_H_ks_level(ev) ...  \  
0            5.70     -5.329     -0.698     3.014      -2.874 ...  
1           95.75     -5.329     -0.698     3.014      -2.874 ...  
2          102.55     -5.329     -0.698     3.014      -2.874 ...  
3          102.55     -5.329     -0.698     3.014      -2.874 ...  
4          101.35     -5.329     -0.698     3.014      -2.874 ...  
..          ...        ...          ...          ...          ...  ...  
178        101.35     -5.516     0.278      2.619      -3.346 ...  
179        32.00      -5.516     0.278      2.619      -3.346 ...  
180        111.50     -5.516     0.278      2.619      -3.346 ...  
181        76.10      -5.516     0.278      2.619      -3.346 ...
```

182	51.20	-5.516	0.278	2.619		-3.346	...	
	B4_number	B5_NNN	B5_number	B6_NNN	B6_number	B7_NNN	B7_number	\
0	-	-	-	-	-	-	-	-
1	-	-	-	-	-	-	-	-
2	-	-	-	-	-	-	-	-
3	-	-	-	-	-	-	-	-
4	-	-	-	-	-	-	-	-
..
178	-	-	-	-	-	-	-	-
179	-	-	-	-	-	-	-	-
180	-	-	-	-	-	-	-	-
181	-	-	-	-	-	-	-	-
182	-	-	-	-	-	-	-	-
	B8_NNN	B8_number	Vol(A3)					
0	-	-	95.88					
1	-	-	45.10					
2	-	-	98.15					
3	-	-	66.96					
4	-	-	68.62					
..					
178	-	-	245.78					
179	-	-	256.05					
180	-	-	429.84					
181	-	-	407.39					
182	-	-	518.86					

[183 rows x 93 columns]

1.1 Select Features and Clean Data

Choose physical and atomic descriptors as features and H(J/mol-atomes) as the target (formation energy).

```
[15]: target = "H(J/mol-atomes)"
features = [
    'SO J/cal/mol)', 'BDE(kj/mol)', 'SA(J/cal/deg)', 'SB(J/cal/mol)',
    'A_IP(ev)', 'A_EA(ev)', 'A_EN(ev)', 'B_IP(ev)', 'B_EA(ev)', 'B_EN(ev)',
    'Vol(A3)', 'length_a(A)', 'length_b(A)', 'length_c(A)',
    'A_rs(A)', 'A_rp(A)', 'A_rd(A)', 'B_rs(A)', 'B_rp(A)', 'B_rd(A)',
    'Z', 'SG_Number', 'R_Value'
]
df_clean = df[features + [target]].replace("~-", pd.NA)
df_clean = df_clean.dropna()
df_clean = df_clean.astype(float)
```

```

X = df_clean[features]
y = df_clean[target]
X.head(), y.head()

[15]: (   SO J/cal/mol)  BDE(kj/mol)  SA(J/cal/deg)  SB(J/cal/mol)  A_IP(ev)  \
0      5.770000          0.0          29.1          5.70       -5.329
1      4.681075          0.0          29.1         95.75       -5.329
2      5.298567         340.5          29.1        102.55       -5.329
3      5.770000         340.5          29.1        102.55       -5.329
4      5.770000         577.0          29.1        101.35       -5.329

    A_EA(ev)  A_EN(ev)  B_IP(ev)  B_EA(ev)  B_EN(ev)  ...  length_c(A)  \
0     -0.698     3.014   -10.852    -0.872     5.862  ...      5.43440
1     -0.698     3.014   -13.585    -1.867     7.726  ...      3.87932
2     -0.698     3.014   -16.433    -3.006     9.720  ...      4.61280
3     -0.698     3.014   -16.433    -3.006     9.720  ...      7.72400
4     -0.698     3.014   -19.404    -4.273    11.839  ...      4.09400

    A_rs(A)  A_rp(A)  A_rd(A)  B_rs(A)  B_rp(A)  B_rd(A)    Z  SG_Number  \
0     1.652     1.995     6.93     0.644     0.630     1.631   2.0      71.0
1     1.652     1.995     6.93     0.539     0.511     1.540   1.0     191.0
2     1.652     1.995     6.93     0.462     0.427     2.219   4.0     225.0
3     1.652     1.995     6.93     0.462     0.427     2.219   2.0     194.0
4     1.652     1.995     6.93     0.406     0.371     1.428   4.0     225.0

    R_Value
0     0.0210
1     0.0000
2     0.0794
3     0.0000
4     0.0000

[5 rows x 23 columns],
0     -14850.0
1     -41200.0
2    -199300.0
3    -158475.0
4    -307500.0
Name: H(J/mol-atomes), dtype: float64)

```

1.2 Train Model

Use Random Forest Regressor model

```
[16]: from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_absolute_error, r2_score
```

```

# Split the data & initialize & train
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
model = RandomForestRegressor(n_estimators=100, random_state=42)
model.fit(X_train, y_train)

# Predict on test data & evaluate
y_pred = model.predict(X_test)
mae = mean_absolute_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print(f"Mean Absolute Error is: {mae:.2f}")
print(f"R2 Score is: {r2:.2f}")

```

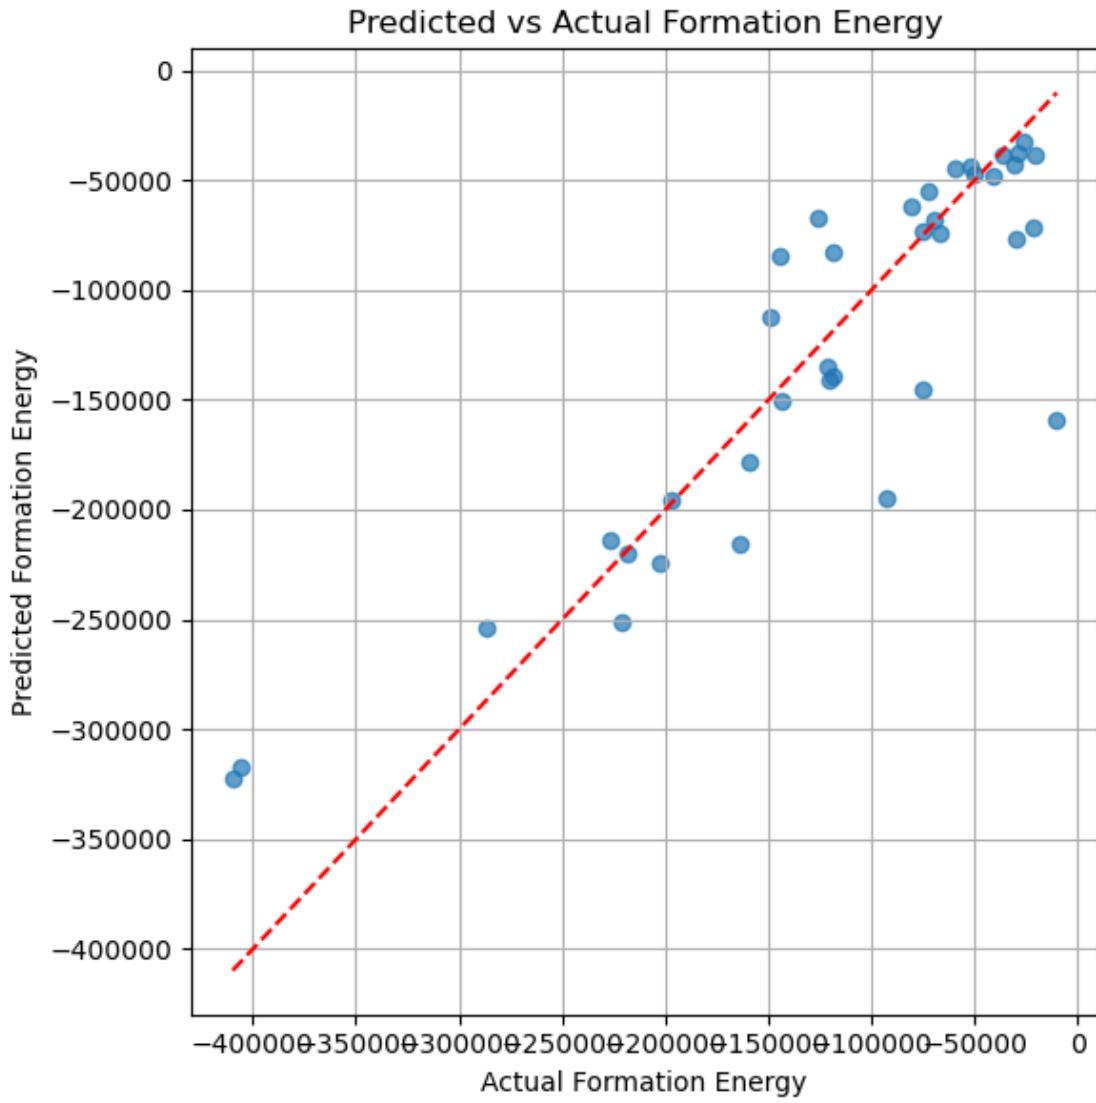
Mean Absolute Error is: 30972.76

R² Score is: 0.78

1.3 Visualization

Plot the predicted vs. actual formation energies and examine features.

```
[17]: import matplotlib.pyplot as plt
plt.figure(figsize=(6, 6))
plt.scatter(y_test, y_pred, alpha=0.7)
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'r--')
plt.xlabel("Actual Formation Energy")
plt.ylabel("Predicted Formation Energy")
plt.title("Predicted vs Actual Formation Energy")
plt.grid(True)
plt.tight_layout()
plt.show()
```



```
[18]: importances = model.feature_importances_
feature_importance_df = pd.DataFrame({
    'Feature': X.columns,
    'Importance': importances
}).sort_values(by='Importance', ascending=False)

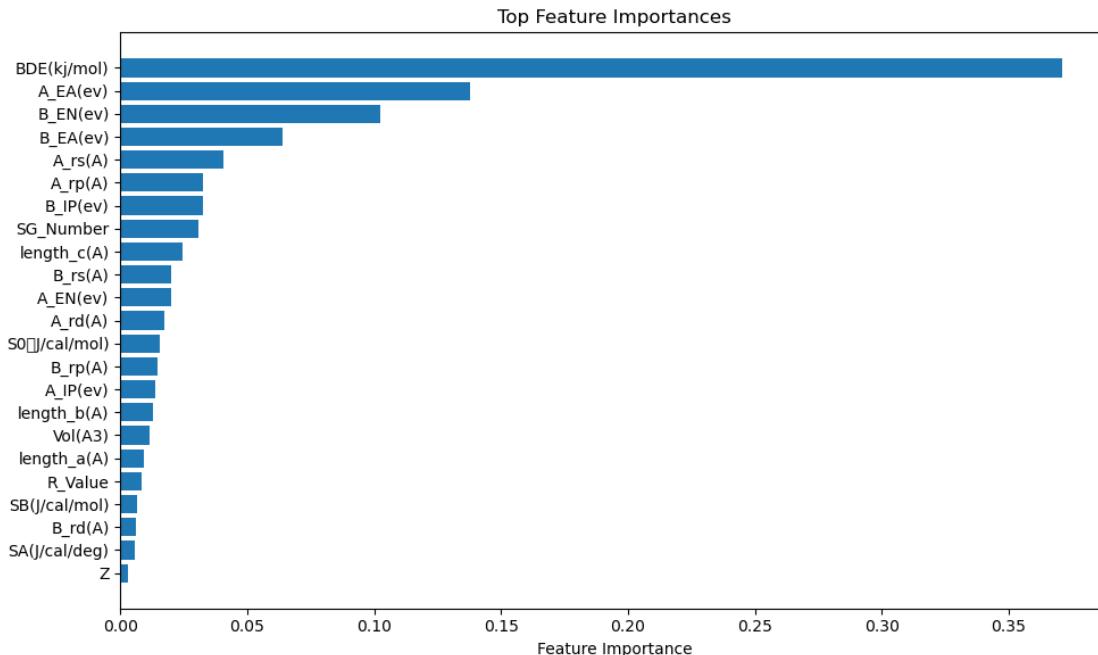
# Bar plot of top features
plt.figure(figsize=(10, 6))
plt.barh(feature_importance_df['Feature'], feature_importance_df['Importance'])
plt.xlabel("Feature Importance")
plt.title("Top Feature Importances")
plt.gca().invert_yaxis()
plt.tight_layout()
```

```

plt.show()

/tmp/ipykernel_686/1796090255.py:13: UserWarning: Glyph 65288 (\N{FULLWIDTH LEFT PARENTHESIS}) missing from current font.
    plt.tight_layout()
/opt/conda/lib/python3.12/site-packages/IPython/core/pylabtools.py:170:
UserWarning: Glyph 65288 (\N{FULLWIDTH LEFT PARENTHESIS}) missing from current font.
fig.canvas.print_figure(bytes_io, **kw)

```



1.4 Model Comparison

Compare our Random Forest model to a Gradient Boosting Regressor.

```
[19]: from sklearn.ensemble import GradientBoostingRegressor
gbr = GradientBoostingRegressor(n_estimators=100, random_state=42)
gbr.fit(X_train, y_train)
y_pred_gbr = gbr.predict(X_test)
mae_gbr = mean_absolute_error(y_test, y_pred_gbr)
r2_gbr = r2_score(y_test, y_pred_gbr)
print(f"MAE: {mae_gbr:.2f}")
print(f"R² Score: {r2_gbr:.2f}")
```

MAE: 28059.17

R² Score: 0.86

```
[20]: from sklearn.model_selection import cross_val_score, cross_val_predict
from sklearn.metrics import mean_absolute_error, make_scorer

# Cross Validation
mae_scorer = make_scorer(mean_absolute_error)
rf_cv_r2 = cross_val_score(model, X, y, cv=5, scoring='r2')
rf_cv_mae = cross_val_score(model, X, y, cv=5, scoring=mae_scorer)

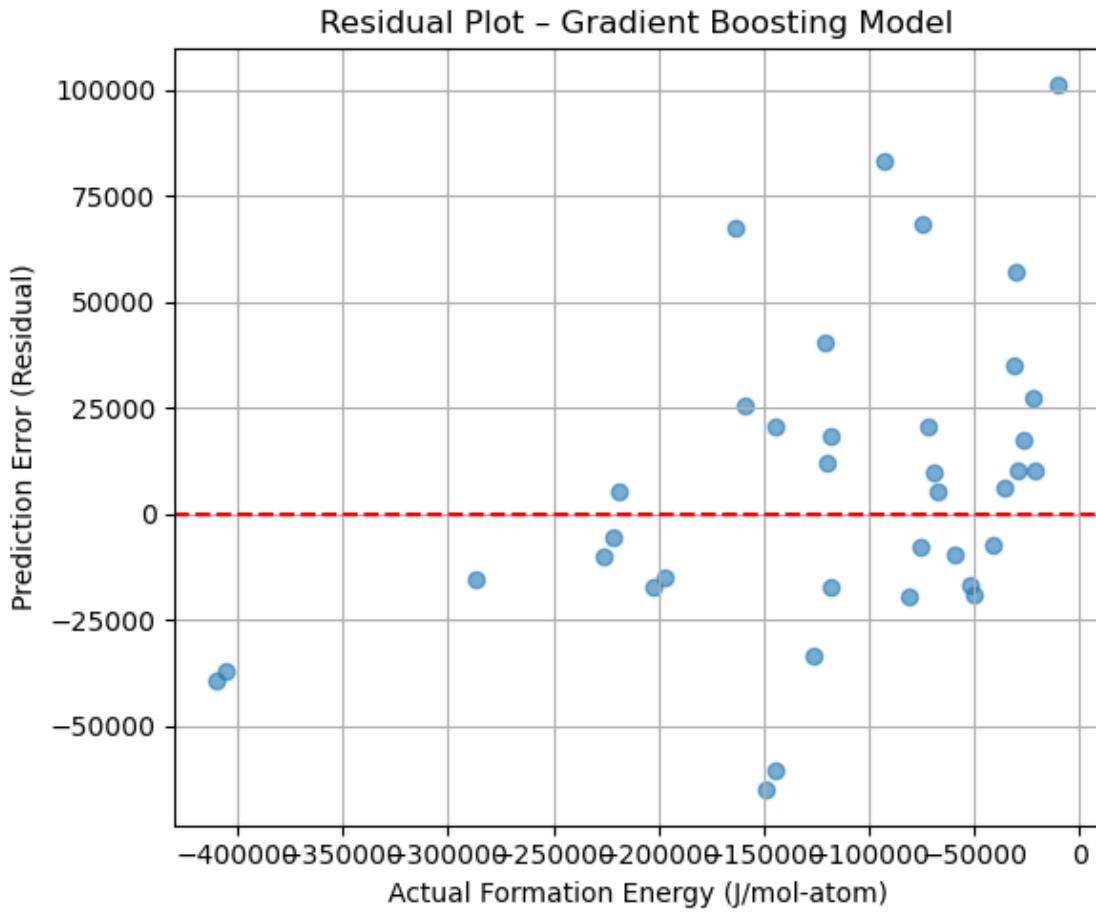
gbr_cv_r2 = cross_val_score(gbr, X, y, cv=5, scoring='r2')
gbr_cv_mae = cross_val_score(gbr, X, y, cv=5, scoring=mae_scorer)

print("Random Forest - CV R2:", rf_cv_r2.mean(), "CV MAE:", rf_cv_mae.mean())
print("Gradient Boosting - CV R2:", gbr_cv_r2.mean(), "CV MAE:", gbr_cv_mae.
     mean())
```

Random Forest - CV R²: 0.7618753893096278 CV MAE: 31896.04232084183
 Gradient Boosting - CV R²: 0.7855547850181079 CV MAE: 29402.273959916365

```
[21]: # Residual plot for best model
residuals = y_test - y_pred_gbr

plt.figure(figsize=(6, 5))
plt.scatter(y_test, residuals, alpha=0.6)
plt.axhline(0, color='red', linestyle='--')
plt.xlabel("Actual Formation Energy (J/mol-atom)")
plt.ylabel("Prediction Error (Residual)")
plt.title("Residual Plot - Gradient Boosting Model")
plt.grid(True)
plt.tight_layout()
plt.show()
```



1.5 Model with Top 5 features

Smaller model with top 5 features and PCA

```
[22]: top_features = feature_importance_df.sort_values(by='Importance', 
    ↴ascending=False)['Feature'].head(5).tolist()
print("Top 5 features:", top_features)
```

Top 5 features: ['BDE(kj/mol)', 'A_EA(ev)', 'B_EN(ev)', 'B_EA(ev)', 'A_rs(A)']

```
[23]: X_small = X[top_features]
X_train_small, X_test_small, y_train_small, y_test_small = 
    ↴train_test_split(X_small, y, test_size=0.2, random_state=42)
gbr_small = GradientBoostingRegressor(n_estimators=100, random_state=42)
gbr_small.fit(X_train_small, y_train_small)
y_pred_small = gbr_small.predict(X_test_small)
mae_small = mean_absolute_error(y_test_small, y_pred_small)
r2_small = r2_score(y_test_small, y_pred_small)
print("Small GBR model - MAE:", mae_small)
```

```
print("Small GBR model - R2:", r2_small)
```

```
Small GBR model - MAE: 29182.571420915418
Small GBR model - R2: 0.8291392845150648
```

```
[24]: from sklearn.decomposition import PCA
from sklearn.pipeline import make_pipeline
from sklearn.preprocessing import StandardScaler
pca_pipeline = make_pipeline(StandardScaler(), PCA(n_components=5), GradientBoostingRegressor(n_estimators=100, random_state=42))
pca_pipeline.fit(X_train, y_train)
y_pred_pca = pca_pipeline.predict(X_test)
mae_pca = mean_absolute_error(y_test, y_pred_pca)
r2_pca = r2_score(y_test, y_pred_pca)

print("PCA + GBR model - MAE:", mae_pca)
print("PCA + GBR model - R2:", r2_pca)
```

```
PCA + GBR model - MAE: 42237.20621168072
PCA + GBR model - R2: 0.652603372243924
```

1.5.1 Summary

The goal of this project was to develop a machine learning model that could accurately predict the formation energy of binary compounds using only atomic and compositional features to support faster screening of materials.

The Gradient Boosting model showed superior performance compared to Random Forest model, achieving an R² score of 0.86 and a mean absolute error (MAE) of 28,000 J/mol-atom, which suggests that the model is able to make fairly accurate predictions. To further evaluate model reliability, I applied 5-fold cross-validation in which Gradient Boosting remained the best performer with a cross-validated R² of 0.79 and MAE of 29,400 J/mol-atom.